




# Flink Training for Real-Time Data Engineering

The Shift to Real-Time: From Batch to Streaming

---

# About Instructor

- About Instructor ... *<text size should be 16 and style should be Trebuchet MS>*



**NAME**

Datacouch Instructor



# AGENDA

- Introduction to Apache Flink
- Overview of Flink's architecture



# What is Apache Flink?

**Apache Flink** is an open-source, distributed engine for stateful processing over unbounded (streams) and bounded (batches) data sets.

- Runs continuously with minimal downtime, processing data as ingested.
- Designed for **low latency, in-memory computations**, and **high availability**.
- Removes single points of failure and **scales horizontally**.
- Provides **exactly-once state management**.
- Supports **event-time processing**, handling out-of-order and late data.
- **Streaming-first** framework with a unified API for stream and batch processing.



# Why Use Apache Flink?

Versatile framework for both streaming and batch applications.

## Event-driven applications:

- Ingest events from one or more streams.
- Perform computations, update state, or trigger external actions.
- Supports **stateful processing**, enabling logic that depends on event history.

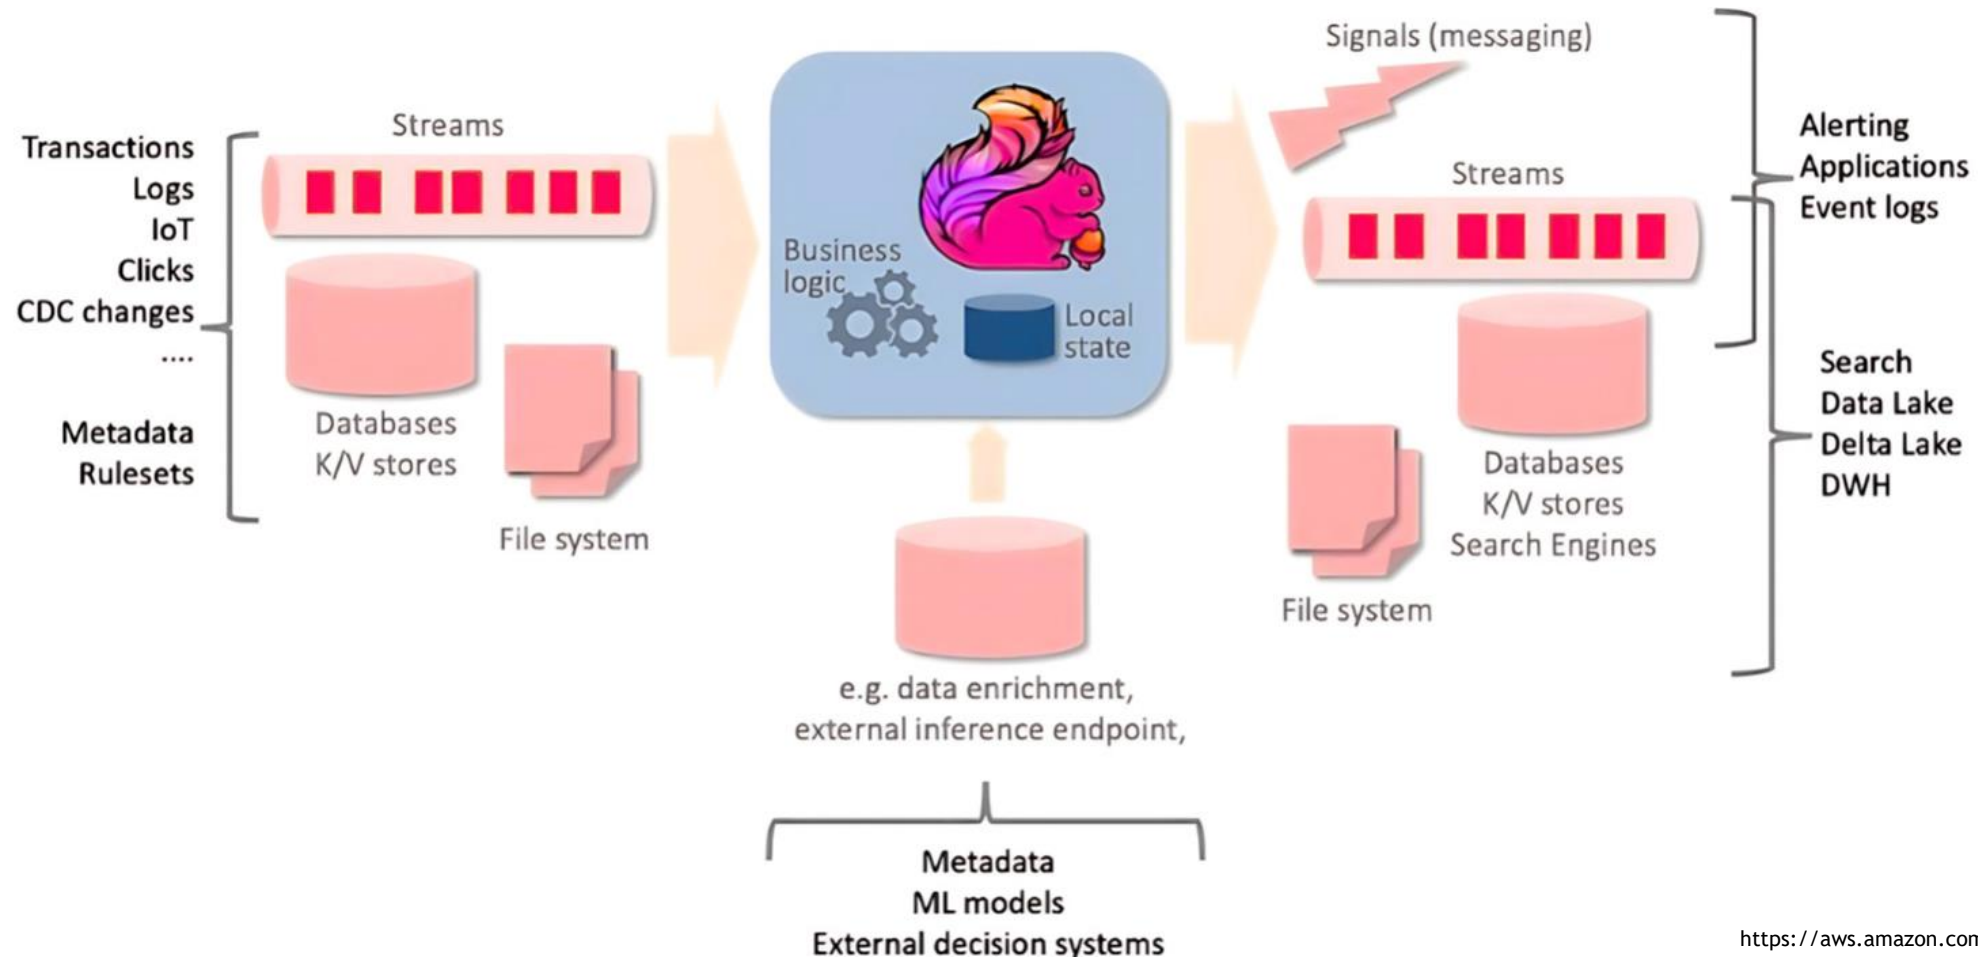
## Data analytics applications:

- Extract insights continuously instead of re-running queries on finite datasets.
- Enable **real-time streaming queries** that emit and update results continuously.

## Data pipeline applications:

- Transform and enrich data while moving it between storages.
- Replace periodic batch ETL with **continuous, low-latency data movement**.

# Why Use Apache Flink?



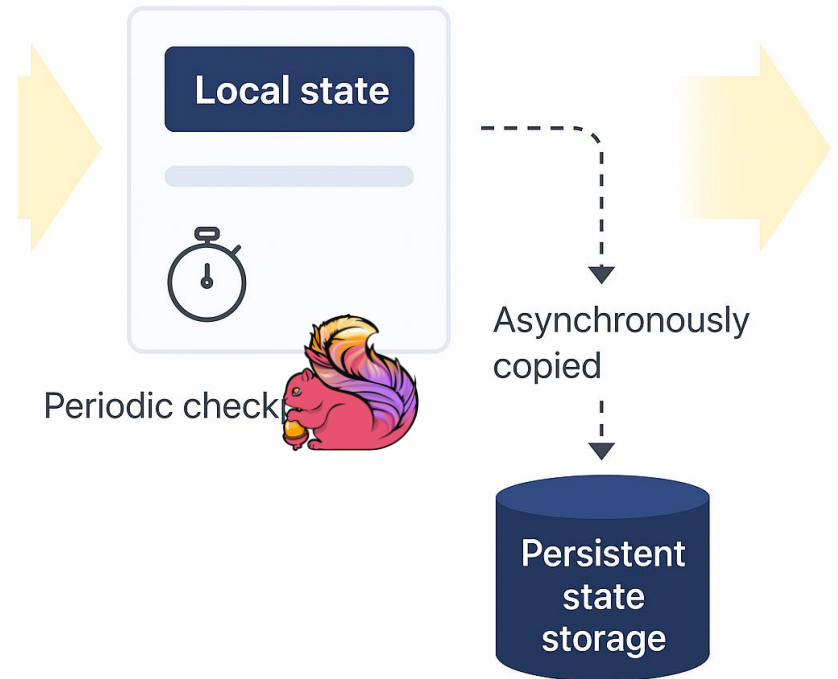
<https://aws.amazon.com/what-is/apache-flink/>



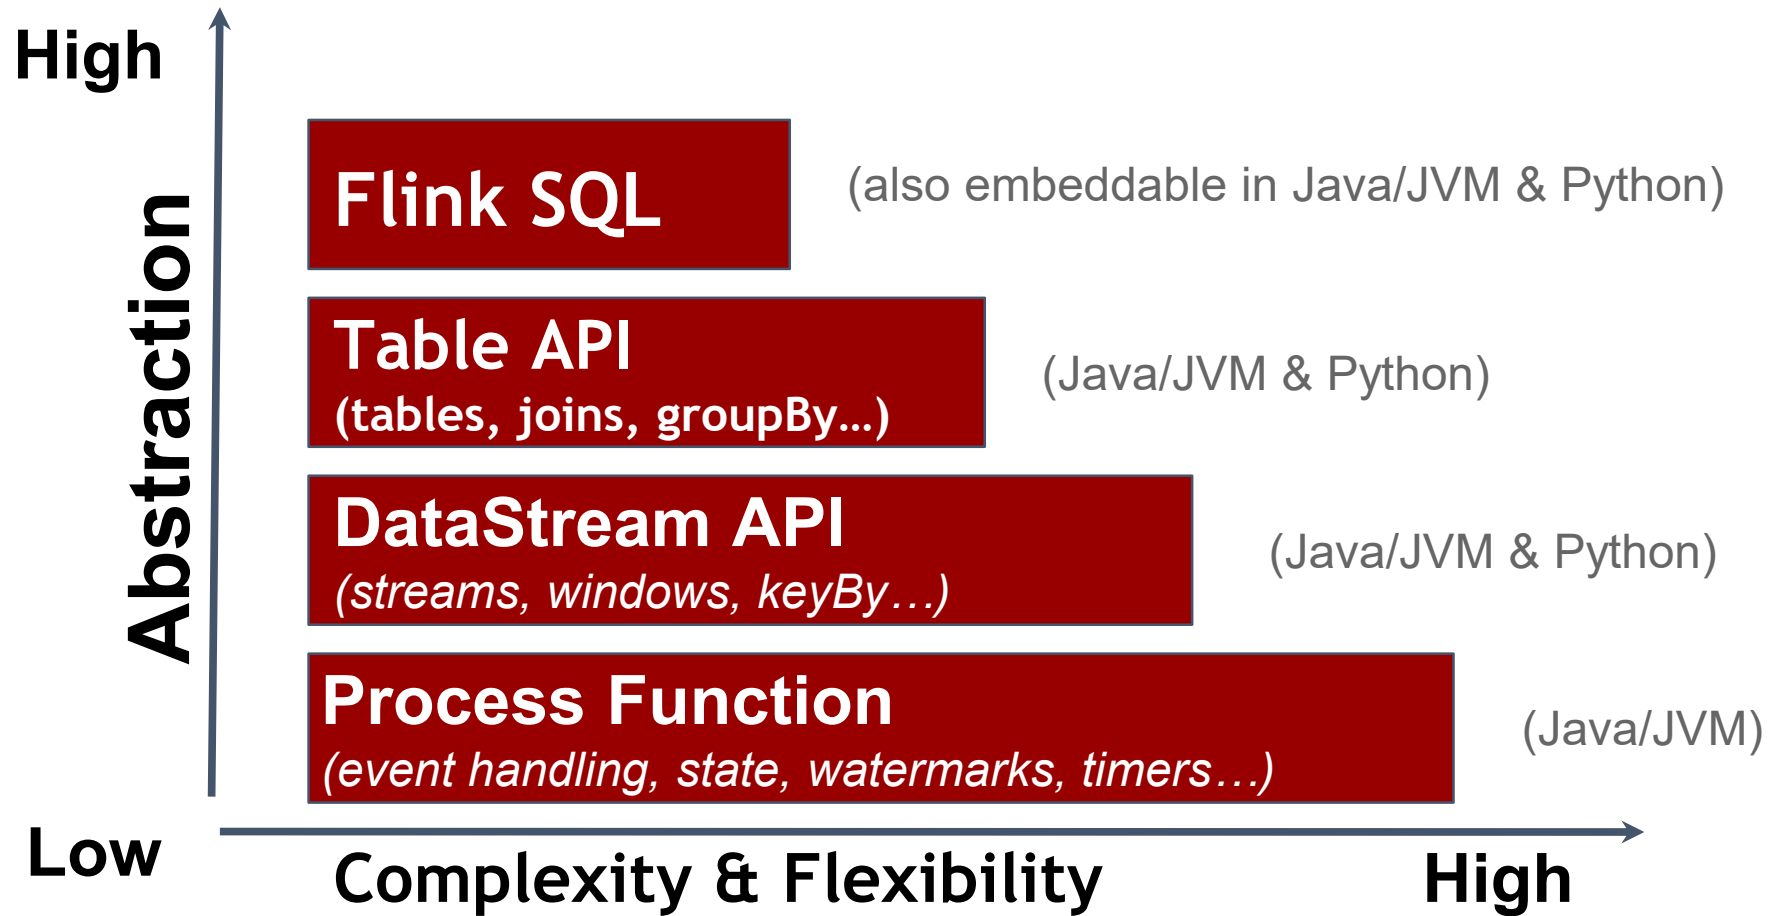
# How does Apache Flink work?

## 1. Dataflow Model

- A Flink application is built as a **dataflow graph**.
  - **Sources**: where data comes in (files, message queues like Kafka, databases, search engines).
  - **Transformations**: operations applied on data streams (e.g., filtering, aggregations, pattern detection).
  - **Sinks**: where results are written (datastores, dashboards, files, etc.).
- The graph is **acyclic** but can be highly complex.
- Processing happens in **real time** with **high throughput** and **low latency**.



# Apache Flink: Multiple Abstractions & Languages





# Why Flink Excels at Stateful Stream Processing

**True Stream-first Engine** - Batch is treated as bounded streams.

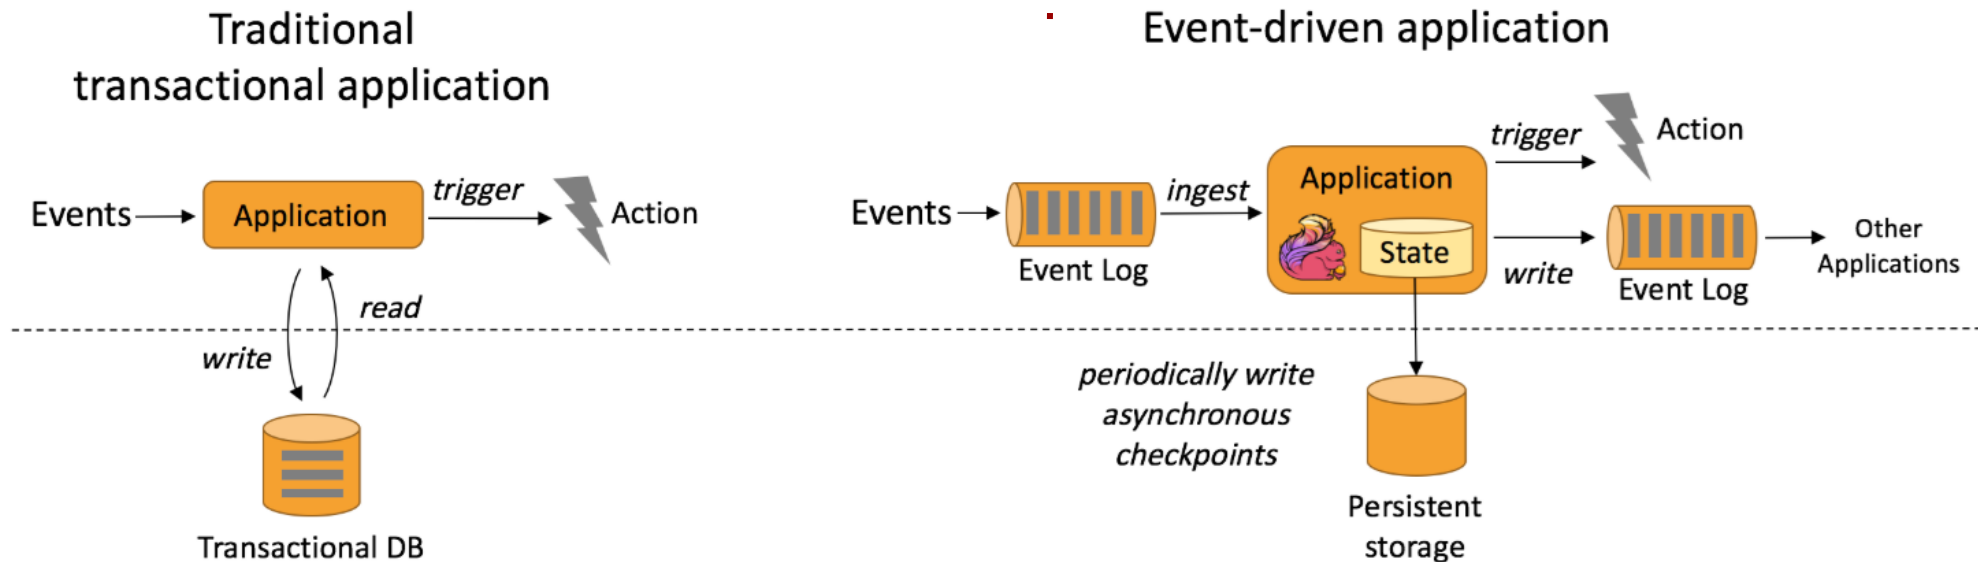
**Stateful Processing** - Applications keep local state for fast access.

**Exactly-once Consistency** - Through checkpoints & savepoints.

**Low Latency + High Throughput** - Designed for real-time workloads.

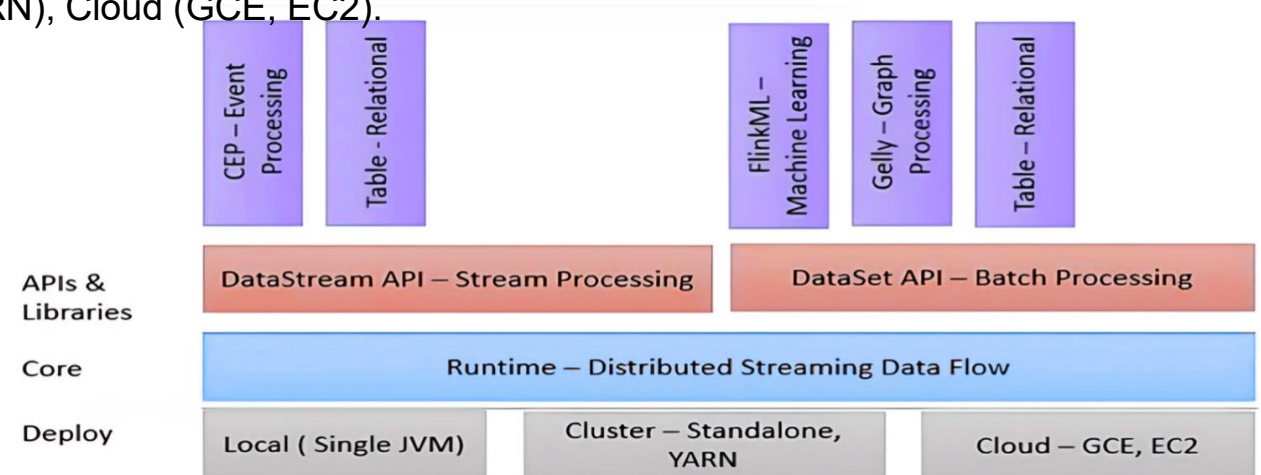
**Scalability & Fault Tolerance** - Distributed, resilient to failures.

**Flexible Deployments** - Local, cluster (YARN, Kubernetes), cloud.



# Flink's architecture

- **APIs & Libraries:** CEP (events), Table/SQL, FlinkML (ML), Gelly (graphs).
- **Core APIs:**
  - **DataStream API** – Stream processing (unbounded).
  - **DataSet API** – Batch processing (bounded).
- **Runtime:** Distributed streaming data flow engine with parallelism & fault tolerance.
- **Deployment:** Local (JVM), Cluster (Standalone/YARN), Cloud (GCE, EC2).



# Jobs, Tasks, and Operators

## Jobs

- Represent a **user-submitted application**.
- Defined as a **JobGraph** (logical graph).
- Built using **sources** → **transformations** → **sinks**.

## Operators

- Fundamental operations: **map, filter, keyBy, reduce, etc.**
- Connected into a **directed acyclic graph (DAG)**.
- Define the **flow of data** in the job.

## Tasks

- **Parallel units of execution** derived from operators.
- Each processes a **data partition independently**.
- Mapped to **available cluster resources**.

# Apache Flink Cluster Model

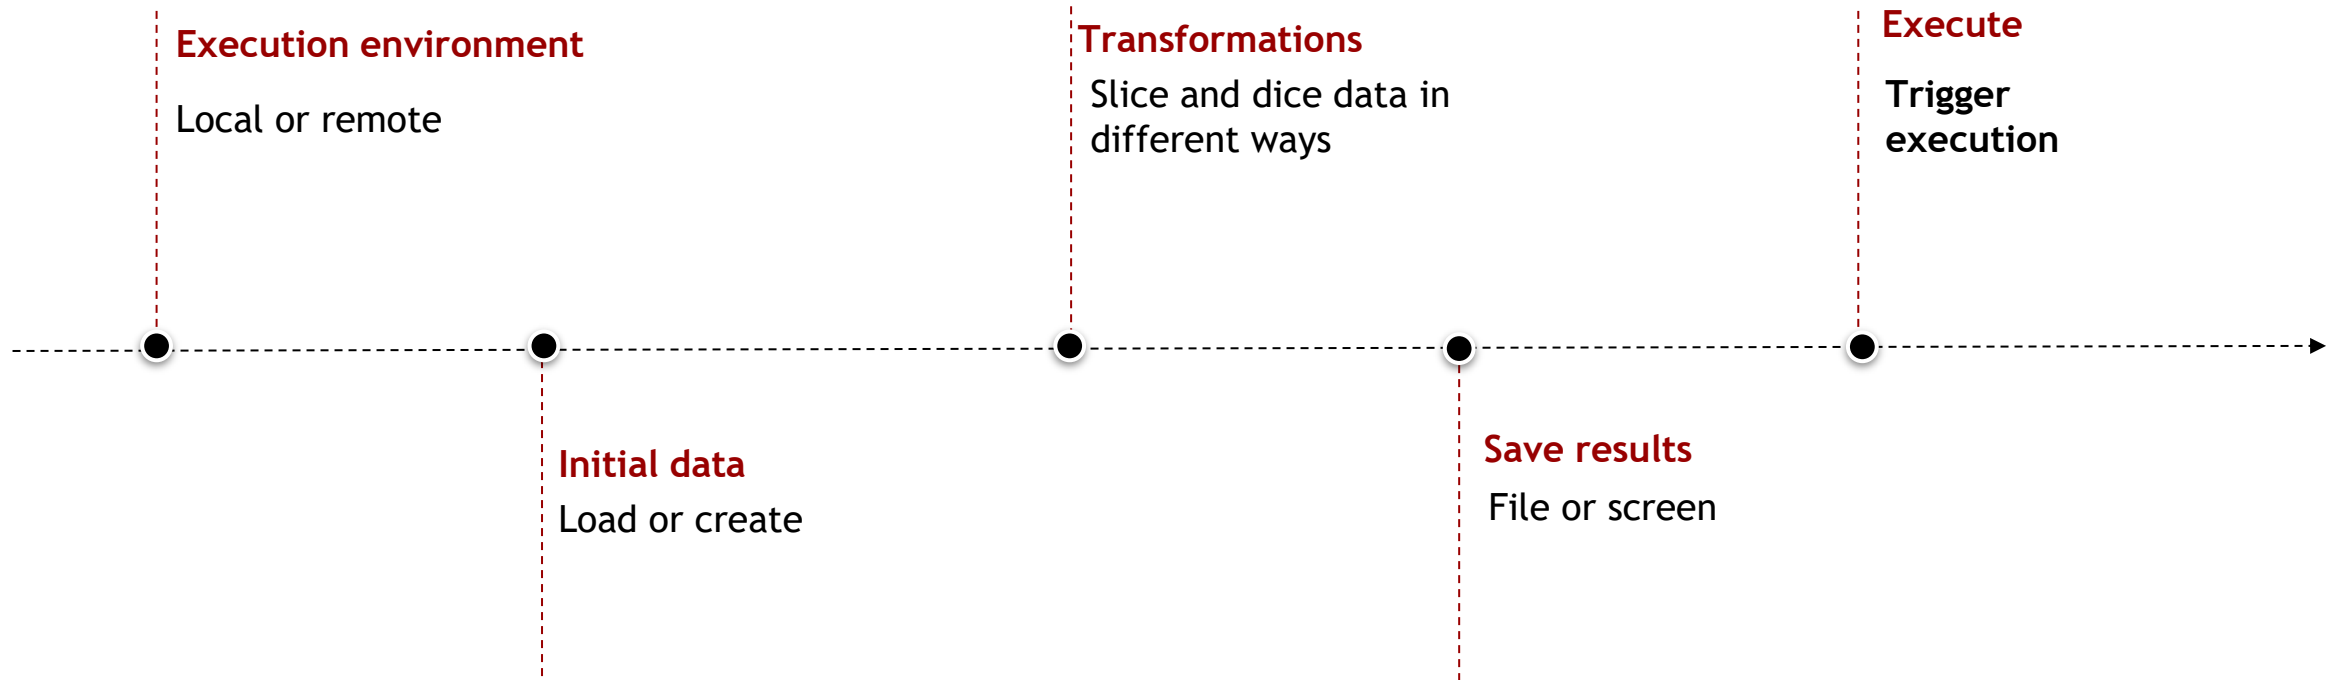
## Master Components

- **Dispatcher**
  - Accepts job submissions (REST, CLI, Web UI).
  - Assigns unique Job IDs & forwards to JobManager.
- **JobManager**
  - Orchestrates job scheduling, checkpointing, state management, recovery.
  - Converts **JobGraph** → **Execution Graph**.
  - Coordinates execution & failover handling.
- **ResourceManager**
  - Integrates with **YARN, Kubernetes, AWS**.
  - Allocates resources & assigns task slots.
  - Enables **elastic scaling**.

## Worker Component

- **TaskManagers**
  - Worker nodes running **slots** for task execution.
  - Execute operators & maintain **local + checkpointed state**.
  - Handle distributed data processing.

# Basic Anatomy of a Flink Program







A low-angle, upward-looking photograph of several modern skyscrapers with glass facades. The image is overlaid with a semi-transparent dark grey rectangle in the center, which contains the text 'THANK YOU'. Additionally, there are three solid red rectangular blocks: one at the top center, one at the bottom left, and one at the bottom right. A small black horizontal bar is positioned below the red block at the bottom left.

**THANK YOU**