



# Flink Training for Real-Time Data Engineering

Performance Tuning and Engineering Best Practices for Apache Flink

---



# AGENDA

- Monitoring Flink jobs: Identifying backpressure and bottlenecks.
- Resource management: Allocating CPU, memory, and parallelism effectively.
- Data Serialization and Schema Management (e.g., using Avro).
- Structuring Flink projects for maintainability and testing.

# Why Performance Tuning Matters in Flink

## Cost Reduction

Pinterest cut streaming costs by 40% and onboarded 40% more jobs through effective tuning strategies

## System Stability

Poor tuning causes backpressure, bottlenecks, and unstable clusters that affect entire data pipelines

## Operational Excellence

Efficient tuning balances throughput, latency, and resource usage for real-time streaming success

# Monitoring Flink Jobs: Spotting Backpressure & Bottlenecks

## Understanding Backpressure

**Backpressure** occurs when CPU starvation or network delays cause task slowdowns throughout your streaming pipeline.

Use Flink UI's **Metrics & Checkpoints** tab to identify heavy operators and lagging subtasks effectively.

### Key Warning Indicators:

- Rising checkpoint durations
- Skewed task load distribution
- Network buffer saturation

# Visualising Backpressure: The Hot Node Phenomenon

## Noisy Neighbours Problem

One job overconsuming CPU resources impacts all other jobs running on the cluster, creating cascading performance issues.

## CPU Starvation Effects

CPU starvation leads to cascading delays and complete pipeline stalls, affecting downstream processing and data freshness.

# Resource Management: CPU, Memory & Parallelism

01

---

## Autoscaling Triggers

Autoscaling activates when CPU exceeds 75% for 15 minutes; manual scaling needed for other resource spikes

02

---

## Operator Parallelism

Set operator parallelism based on resource intensity, not just application-level parallelism settings

03

---

## Task Slot Ratio

Typical stable ratio: 4 operator subtasks per task slot; adjust based on workload intensity requirements

04

---

## KeyBy Partitioning

Use `KeyBy` partitioning to evenly distribute workload and avoid data skew issues

# Best Practices for Parallelism & Scaling

1

## Balance Partitioning

Over-partitioning creates overhead; under-partitioning causes bottlenecks. Find the optimal balance for your workload.

2

## Dynamic Tuning

Tune operator parallelism dynamically via runtime properties without requiring complete redeployment cycles.

3

## Monitor Distribution

Monitor skewed key distributions causing uneven load and backpressure across your cluster nodes.

4

## Task Chaining

Use task chaining to reduce network overhead but watch for resource contention between operators.

# Profiling & Tools for Continuous Performance Insight



## **Async-profiler**

CPU cycles, heap allocations, and flame graphs for precise hotspot detection and performance bottleneck analysis.



## **VisualVM**

Live heap and CPU monitoring capabilities for interactive debugging and real-time performance assessment.



## **jemalloc + jeprof**

Memory profiling to detect leaks over time and understand allocation patterns in long-running applications.



## **Eclipse Memory Analyser**

Deep JVM heap dump analysis for comprehensive memory issue diagnosis and resolution strategies.



# Why Data Serialization & Schema Management Matter



## Performance

Efficient serialization reduces latency and resource consumption in streaming pipelines



## Compatibility

Schema management ensures data compatibility and evolution without system downtime



## Flink Power

Native support for serialization and schema evolution creates robust streaming applications

# Avro & Schema Registry Integration in Flink



## Compact Serialization

Avro provides compact, schema-based binary serialization with strong evolution support



## Centralized Management

Schema Registry centralizes schema storage, versioning, and compatibility enforcement



## Automatic Processing

Flink integrates with Schema Registry for automatic serialization/deserialization of Kafka topics



Register Avro schema via REST API, then use  
`ClouderaRegistryAvroKafkaRecordSerializationSchema` in Flink

# Testing Strategies for Flink Applications

## Unit Testing

Test POJO serialization with Flink's `PojoTestUtils.assertSerializedAsPojo()` for validation

## Integration Testing

Use embedded Kafka and Schema Registry to validate end-to-end serialization workflows

## State Evolution Testing

Employ savepoint-based tests to verify state schema evolution and migration correctness

## CI/CD Automation

Automate schema compatibility checks as part of continuous integration and deployment pipelines

# Ready to Build Robust Flink Pipelines?

## **Schema-Driven Design**

Embrace schema-driven design with Avro and Schema Registry for consistency

## **Clear Structure**

Structure your Flink projects for clarity and comprehensive testability

## **Automated Validation**

Automate serialization and schema validation in your CI/CD processes

## **Production Excellence**

Unlock the full power of Flink's serialization and state management for production-grade streaming



A low-angle, upward-looking photograph of several modern skyscrapers with glass facades. The image is overlaid with a semi-transparent dark grey horizontal band across the middle. Three solid red rectangular blocks are positioned: one at the top center, one at the bottom left, and one at the bottom right. The text 'THANK YOU' is centered within the dark band in a white, bold, sans-serif font.

**THANK YOU**